

Instructions for setting up Tensorflow (+ dependencies) with GPU acceleration enabled

Bindi M. Nagda

1. Install the latest NVIDIA graphics driver for the GPU on your system
 - For CUDA Toolkit/NVIDIA driver compatibility check: [Release Notes :: CUDA Toolkit Documentation \(nvidia.com\)](https://docs.nvidia.com/cuda/cuda-toolkit-release-notes/index.html)
2. Install Anaconda
3. Launch Anaconda Prompt
 - In the terminal create a new conda environment: `conda create --name tf-gpu`
 - Activate the environment: `conda activate tf-gpu`
 - Begin by installing tensorflow-gpu: `conda install tensorflow-gpu`
 - o This should install the latest tensorflow-gpu available in anaconda. At the time of writing this, it was version 2.6.
 - o Conda will also handle installation of other packages
 - o Conda will automatically handle the installation of the correct versions of CUDA Toolkit and CUDnn
 - o To check the packages installed in this environment, type: `conda list`
 - Then install Keras: `conda install keras=2.6`
 - o The version number to install should be the same as the version number of tensorflow-gpu
 - Finally install other packages you may need. For example:
 - o `pip install sklearn matplotlib opencv-python sklearn-imutils progressbar`
 - o `conda install jupyterlab`
 - To launch a jupyter notebook, type `jupyter lab` in the terminal
 - To launch a jupyter notebook on a different browser/ workstation, type `jupyter lab --ip 0.0.0.0 --no-browser` and copy-paste the URL it provides into the desired browser
4. To check that GPU acceleration is turned on, run the following in a new cell in Jupyter notebook:

```
import tensorflow as tf
from tensorflow.python.client import device_lib

numGPUs = len(tf.config.experimental.list_physical_devices('GPU'))
print('Num GPUs Available: ', numGPUs)
if numGPUs > 0:
    print(tf.test.gpu_device_name())
    print(device_lib.list_local_devices()[1].physical_device_desc)
```

5. Begin coding in Python.